



# Database Evolution History

1960s

First Computerized Database Models




1970s

The Dawn of the Database

- The relational model and its language SQL, emerge
- The disruptive model causes the demise of other models

1970 E.F. Codd Writes a Paper on the Relational Database Model



1980s

An Industry Develops

- SQL becomes the de-facto standard
- Commercial offerings from IBM, Oracle grow market
- Other data models enter the scene, without much traction

ORACLE  
1st commercially available RDBMS

IBM DB2

SAP Sybase  
Informix

1990s

Technology Shifts

- Data explodes with the Internet age
- Single server SQL databases run into resource problems
- Business Intelligence and Analytics move out of transactional databases



Teradata

New Distributed SQL Data Warehouse

2000s

New Players Emerge

- Data variety, velocity and volume increase
- New analytics SQL databases are introduced
- NoSQL databases fill the gap for processing unstructured data
- Hadoop gains traction for analyzing petabytes of data

Today

Databases Adapt and Evolve

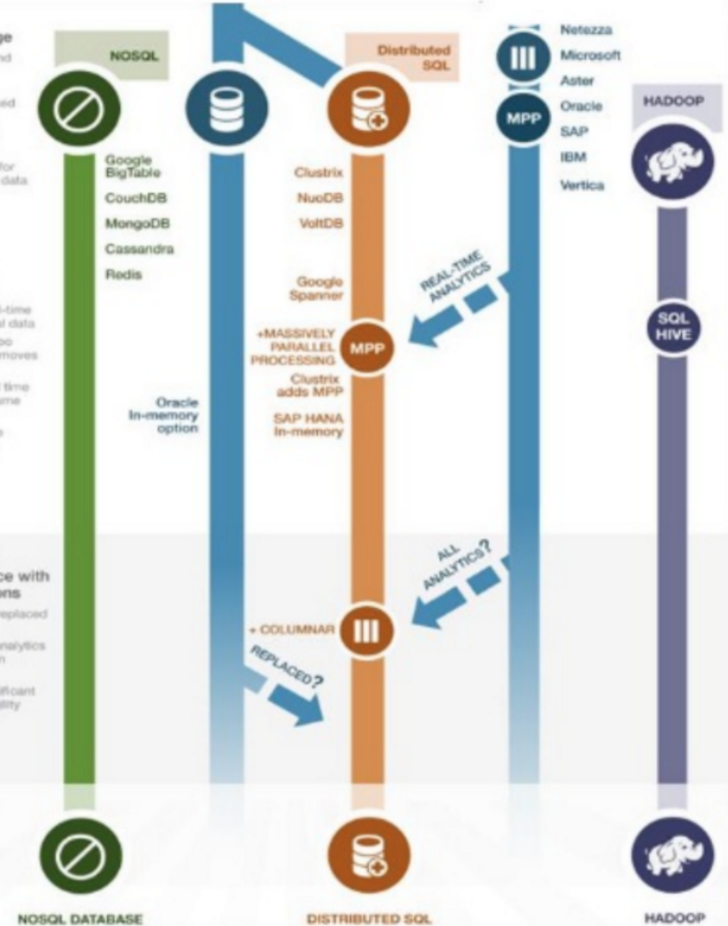
- Businesses require real-time analytics on operational data
- Scale-up SQL proves too costly, but scale-out removes resource constraint
- Scale-out provides real time analytics with high volume transactions
- Google and Clustrix are pioneers in this space

The Future

Businesses Advance with Database Innovations

- Single node SQL gets replaced by scale-out SQL
- Data warehouse type analytics will become available in real-time database
- Businesses gain a significant edge and increased agility

Winning Database Platforms



Source: [Robin Purohit](#)

## 2000s New Players Emerge

- Data variety, velocity and volume increase
- New analytics SQL databases are introduced
- NoSQL databases fill the gap for processing unstructured data
- Hadoop gains traction for analyzing petabytes of data

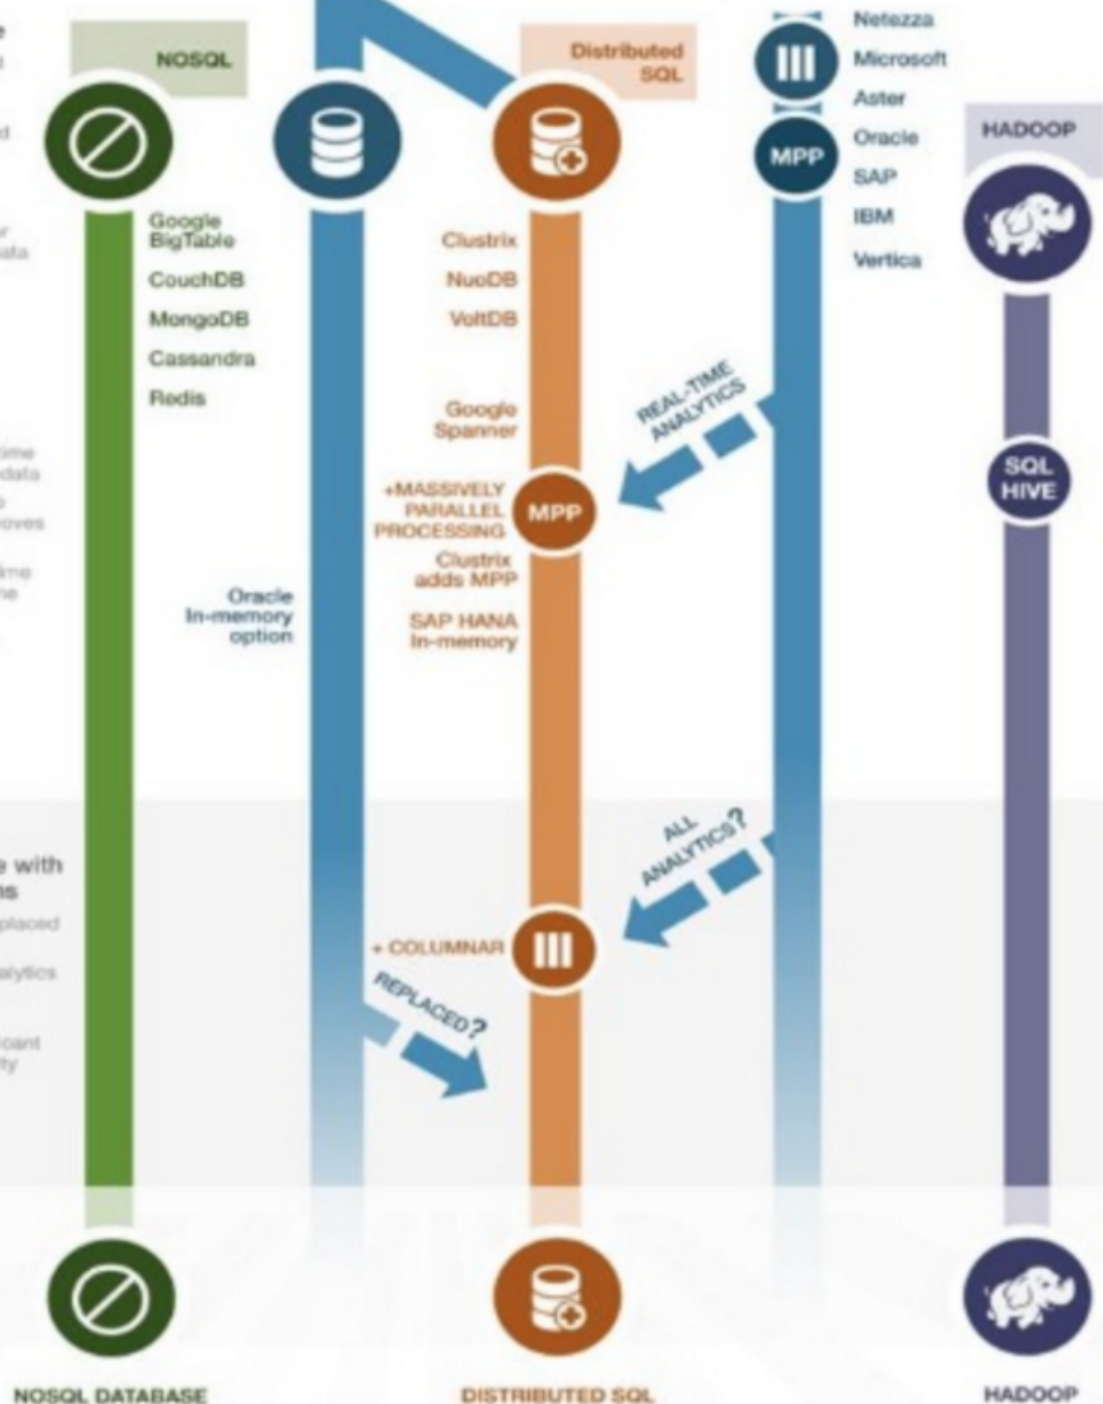
## Today Databases Adapt and Evolve

- Businesses require real-time analytics on operational data
- Scale-up SQL proves too costly, but scale-out removes resource constraint
- Scale-out provides real time analytics with high volume transactions
- Google and Clustrix are pioneers in this space

## The Future Businesses Advance with Database Innovations

- Single node SQL gets replaced by scale-out SQL
- Data warehouse type analytics will become available in real-time database
- Businesses gain a significant edge and increased agility

## Winning Database Platforms



# Technologies du Big-Data

- 16 Nov. Map/Reduce      1CM + 2TDTP
- 23 Nov. Map/Reduce      3TD/TP  
– rendu TP 3 décembre
- 30 Nov. Big-Data & Cognitive (IBM)      2CM  
**présence obligatoire**
- 7 Déc. Big-Data & Cognitive (IBM)      2CM

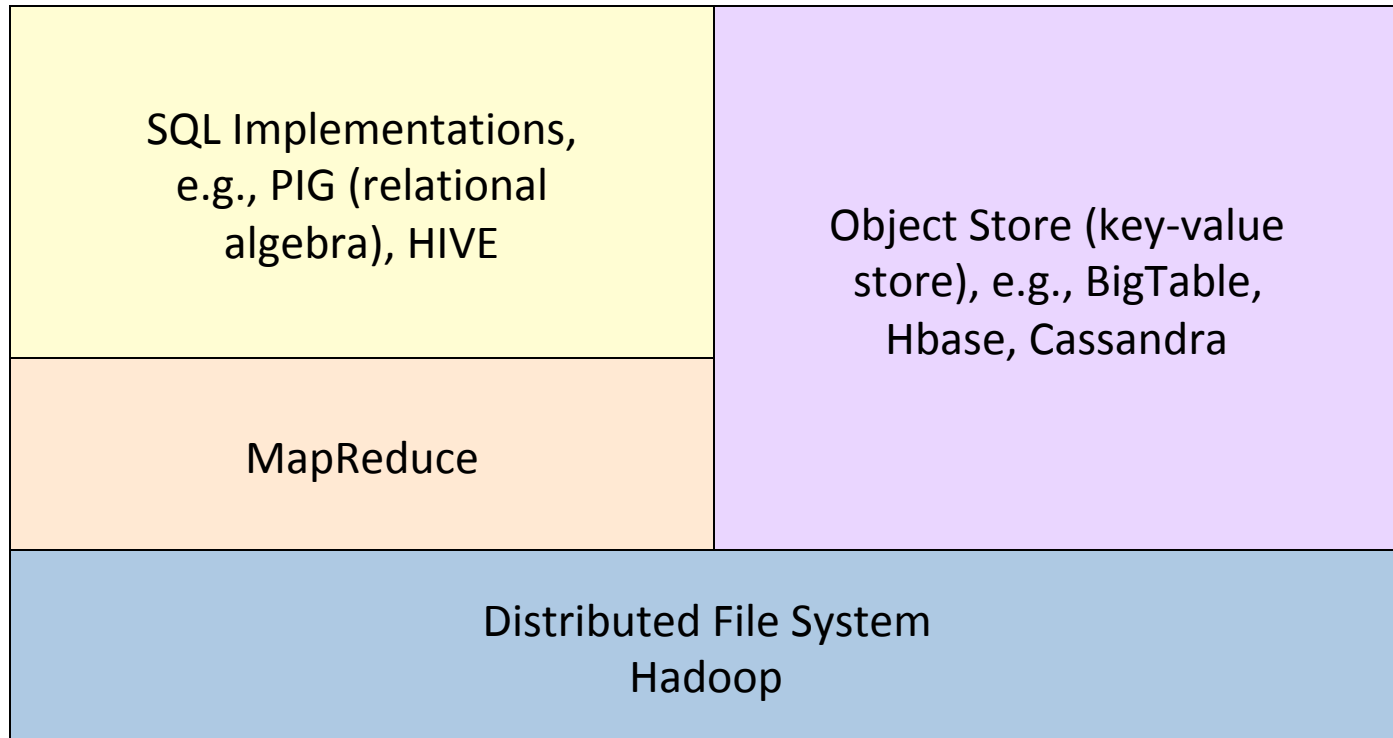
# Map/Reduce

Slides partially collected from  
J. Ullman, J. Leskovec and A.Rajarman

# What is Hadoop-Map/Reduce

- Google's solution to solve Big Data (Volume) problems by means of massive parallelization
- Parallel databases exist since '90, but M/R processing is much more flexible and easier to deploy
- **Hadoop** : distributed file system (holds the data)
- **Map-Reduce** : programming paradigm (compute)

# The New Stack



MAPREDUCE  
(Processing using  
different languages)



HIVE & DRILL  
(Analytical  
SQL-on-Hadoop)



MAHOUT &  
SPARK MLlib  
(Machine learning)



PIG  
(Scripting)



HBASE  
(NoSQL Database)



ZOOKEEPER  
& AMBARI  
(Management  
& Coordination)



SPARK (In-Memory,  
Data Flow Engine)



KAFKA & STORM  
(Streaming)



SOLR & LUCENE  
(Searching  
& Indexing)



OOZIE  
(Scheduling)



Resource  
Management

YARN

Storage



Apache  
Ambari

Source : <http://www.teradata.com.au/Resources/White-Papers/Hadoop-and-the-Data-Warehouse-When-to-Use-Whi>

Requirement	Data Warehouse	Hadoop
Low latency, interactive reports, and OLAP	●	
ANSI 2003 SQL compliance is required	●	
Preprocessing or exploration of raw unstructured data		●
Online archives alternative to tape		●
High-quality cleansed and consistent data	●	
100s to 1000s of concurrent users	●	●*
Discover unknown relationships in the data	●	●
Parallel complex process logic		●
CPU intense analysis	●	●
System, users, and data governance	●	
Many flexible programming languages running in parallel		●
Unrestricted, ungoverned sand box explorations		●
Analysis of provisional data		●
Extensive security and regulatory compliance	●	
Real time data loading and 1 second tactical queries	●	●*



# Hadoop-M/R does three things


1. Makes it easier to run distributed computations
2. Makes it easier to write distribute programs
3. Makes it easier to deal with node failure

# Hadoop-M/R means three things

1. A model of computing
2. A paradigm of programming
3. A distributed-system architecture


# Sequential Computing

- iterate over inputs to compute a given function




# Sequential Computing

- iterate over inputs to compute a given function




# Sequential Computing

- iterate over inputs to compute a given function




# Sequential Computing

- iterate over inputs to compute a given function




# Sequential Computing

- iterate over inputs to compute a given function



# Sequential Computing

- iterate over inputs to compute a given function



note : there is a lot of sequential computing inside parallel programs



# Just before we introduce M/R

- How do you expect it to work ?
- Maybe...
  - read all inputs sequentially
  - then distribute the evaluation of a given function
- Or ...
  - read all inputs in parallel
  - and distribute the evaluation of a given function

# M/R Computing Model

1. **Read** the inputs
2. **Regroup**-the inputs
3. **Evaluate** a function on the regrouped inputs




...and all of these can be done in parallel !

# M/R Computing Model



read




# M/R Computing Model



read

regroup

# M/R Computing Model



read

regroup

evaluate

# M/R Computing Model

Take a set of inputs (files, tables, text...) and a function **F**


1. **Read** (batches of) inputs and assign each input to a group
  - this is called the MAP phase
  - and can be done in parallel *(according to file distribution)*
2. **Regroup** the inputs according to the map criterion :
  - this is called the SHUFFLE phase
  - again, this can be done in parallel *(according to where data is sent)*
3. **Evaluate** the function on the new groups
  - this is called the REDUCE phase
  - again, this can be done in parallel *(according to where data is sent)*

# M/R Computing Model



MAP : read batches  
of inputs


# M/R Computing Model



MAP : read batches  
of inputs




# M/R Computing Model



MAP : read batches  
of inputs

SHUFFLE : regroup  
the inputs





# M/R Computing Model



MAP : read batches  
of inputs

SHUFFLE : regroup  
the inputs





# M/R Computing Model



MAP : read batches  
of inputs

SHUFFLE : regroup  
the inputs

# M/R Computing Model



MAP : read batches  
of inputs

SHUFFLE : regroup  
the inputs

REDUCE : evaluate  
the function

# Text Mining : Keyword Count

Google Chrome - Wikiped x Matthias

en.wikipedia.org/wiki/Google\_Chrome

Create account Log in

Article Talk Read Edit View history Search

## Google Chrome

From Wikipedia, the free encyclopedia

*This article is about the web browser. For the operating system, see Chrome OS.*

**Google Chrome** is a freeware web browser<sup>[10]</sup> developed by Google. It used the WebKit layout engine until version 27 and, with the exception of its iOS releases, from version 28 and beyond uses the WebKit fork Blink.<sup>[11][12][13]</sup> It was first released as a beta version for Microsoft Windows on September 2, 2008, and as a stable public release on December 11, 2008.

As of January 2015, StatCounter estimates that Google Chrome has a 51% worldwide usage share of web browsers, indicating that it is the most widely used web browser in the world.<sup>[14]</sup>


Google releases the majority of Chrome's source code as an open-source project Chromium.<sup>[15][16]</sup> A notable component that is not open source is the built-in Adobe Flash Player.

### Google Chrome




Developer(s)	Google Inc.
Initial release	September 2, 2008; 6 years ago
Stable release	Windows, OS X, Linux 42.0.2311.90 (April 14,

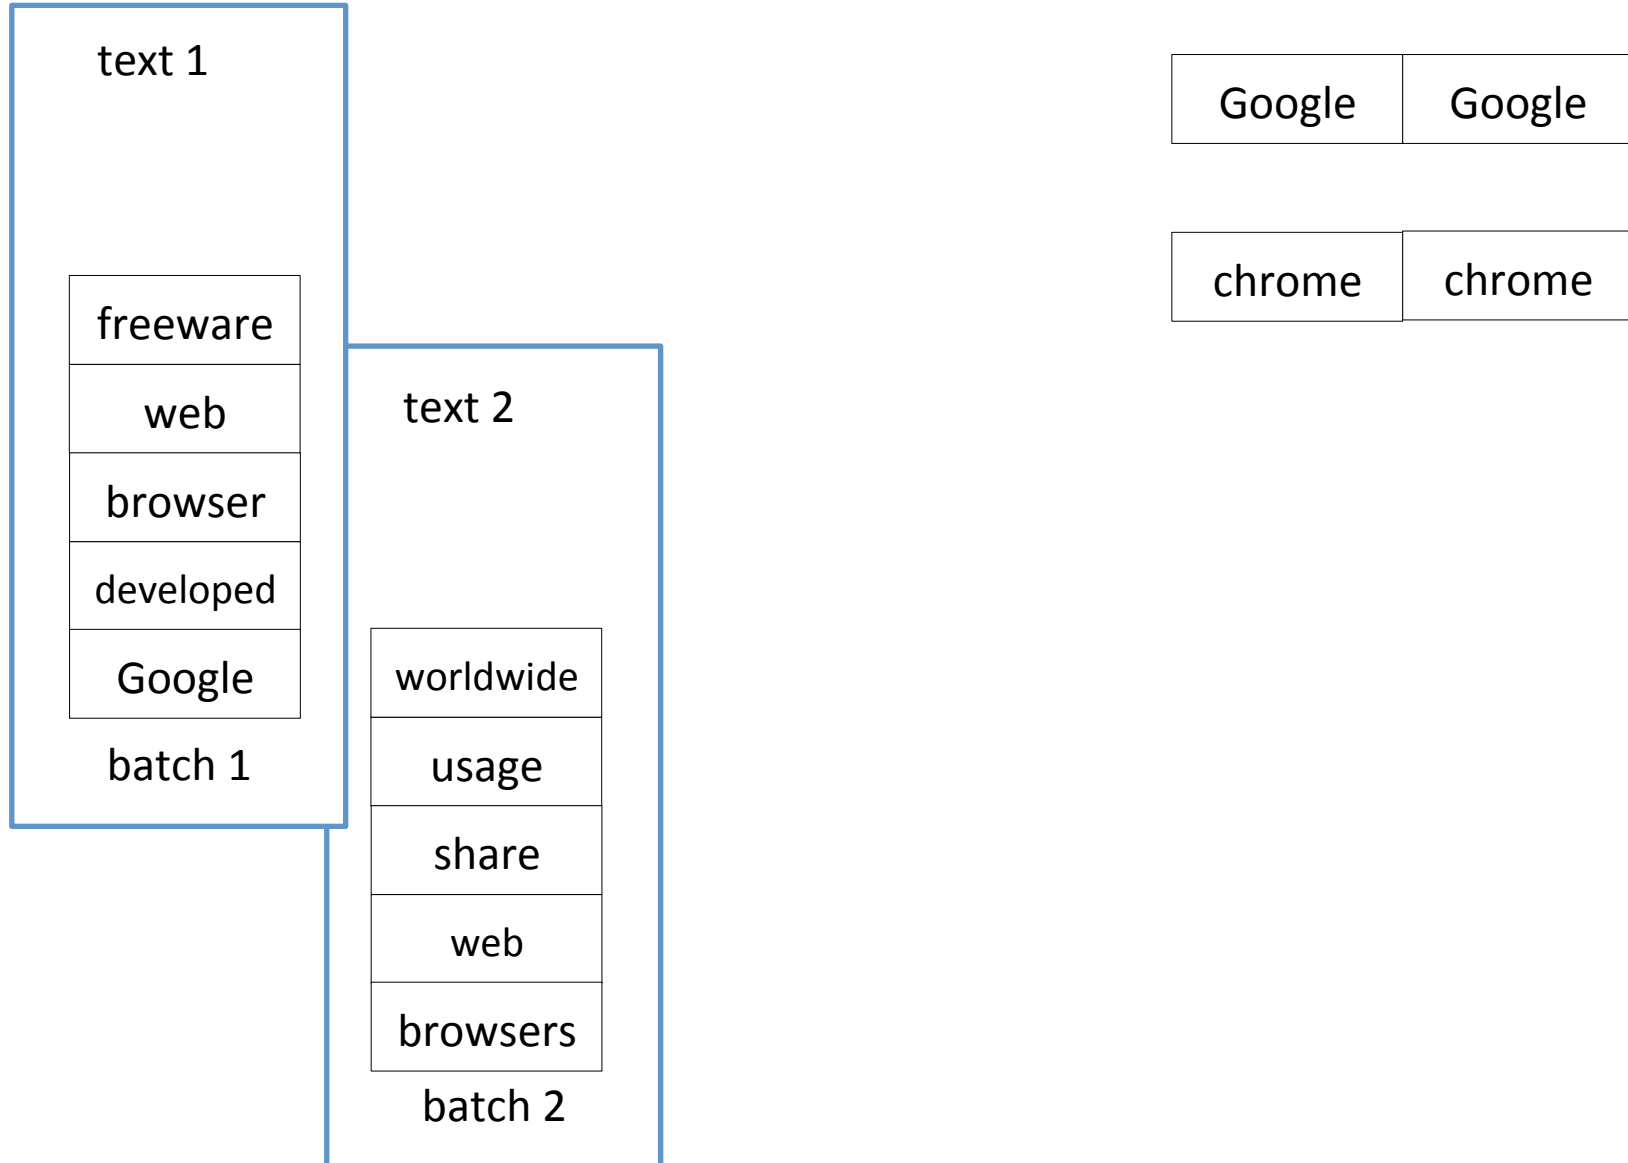
# M/R Computing Model



# M/R Computing Model

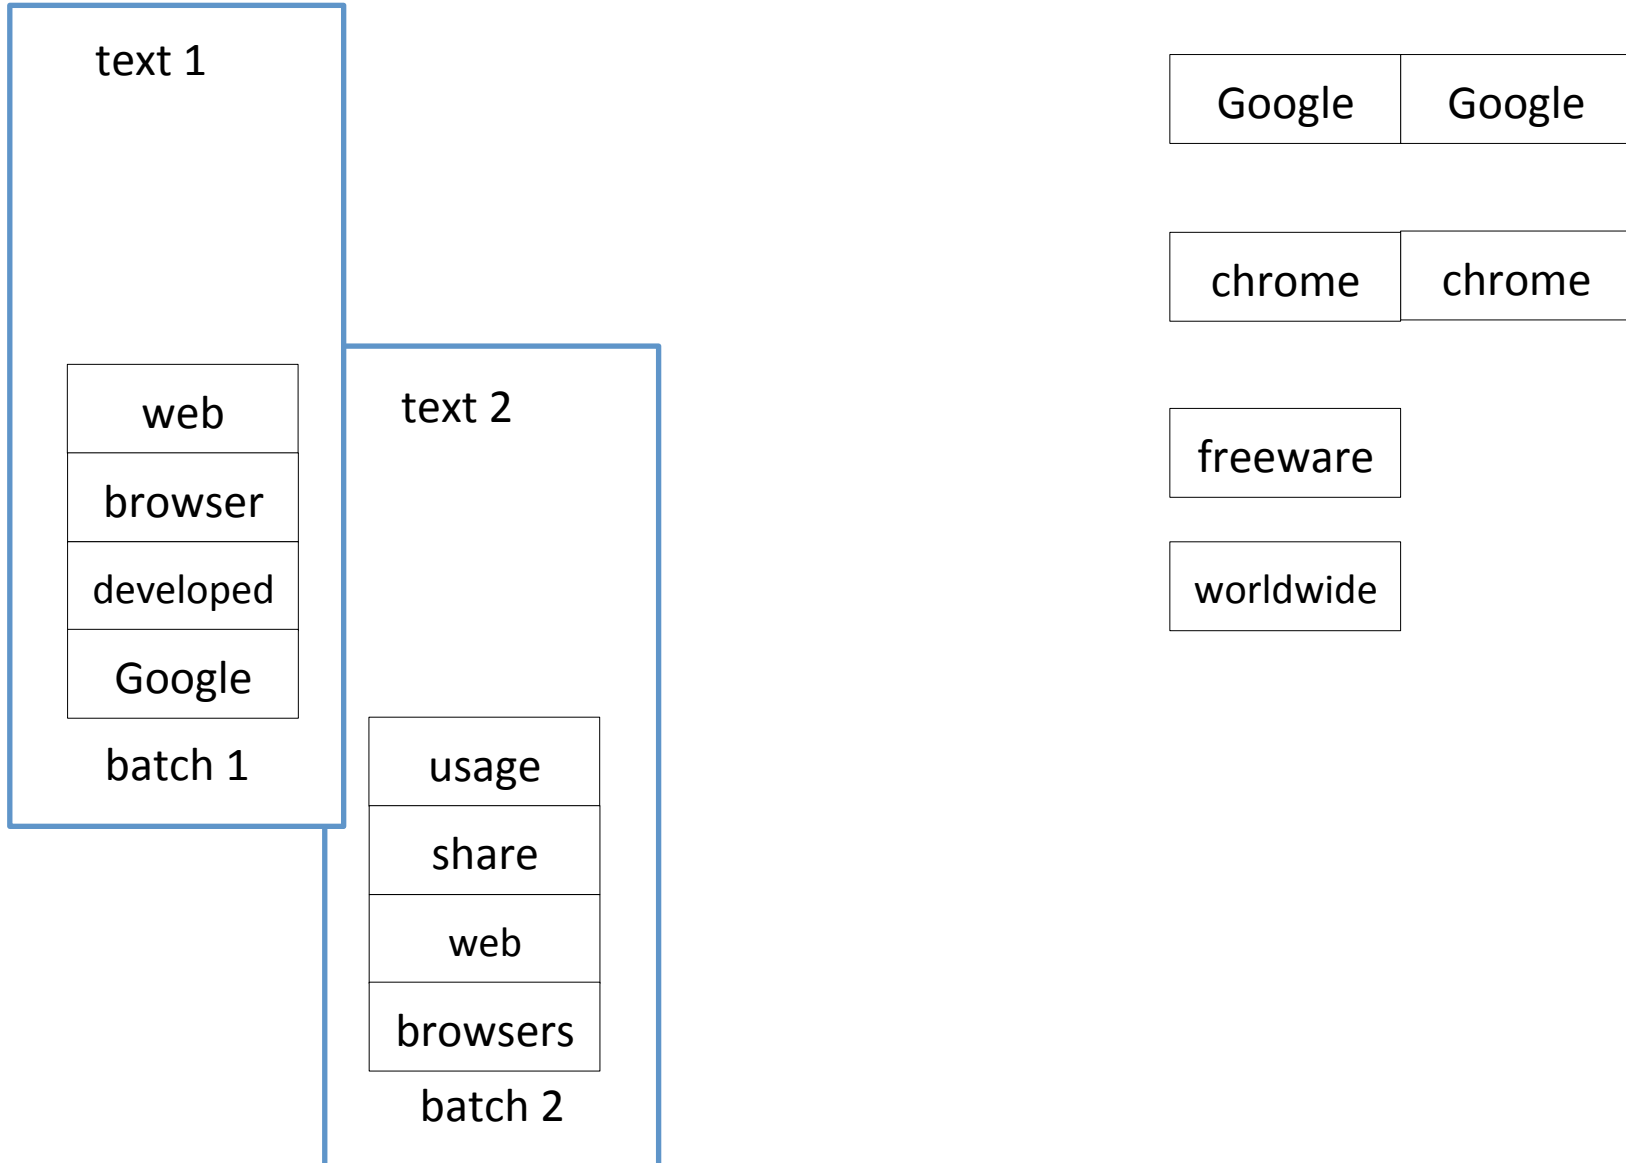


# M/R Computing Model

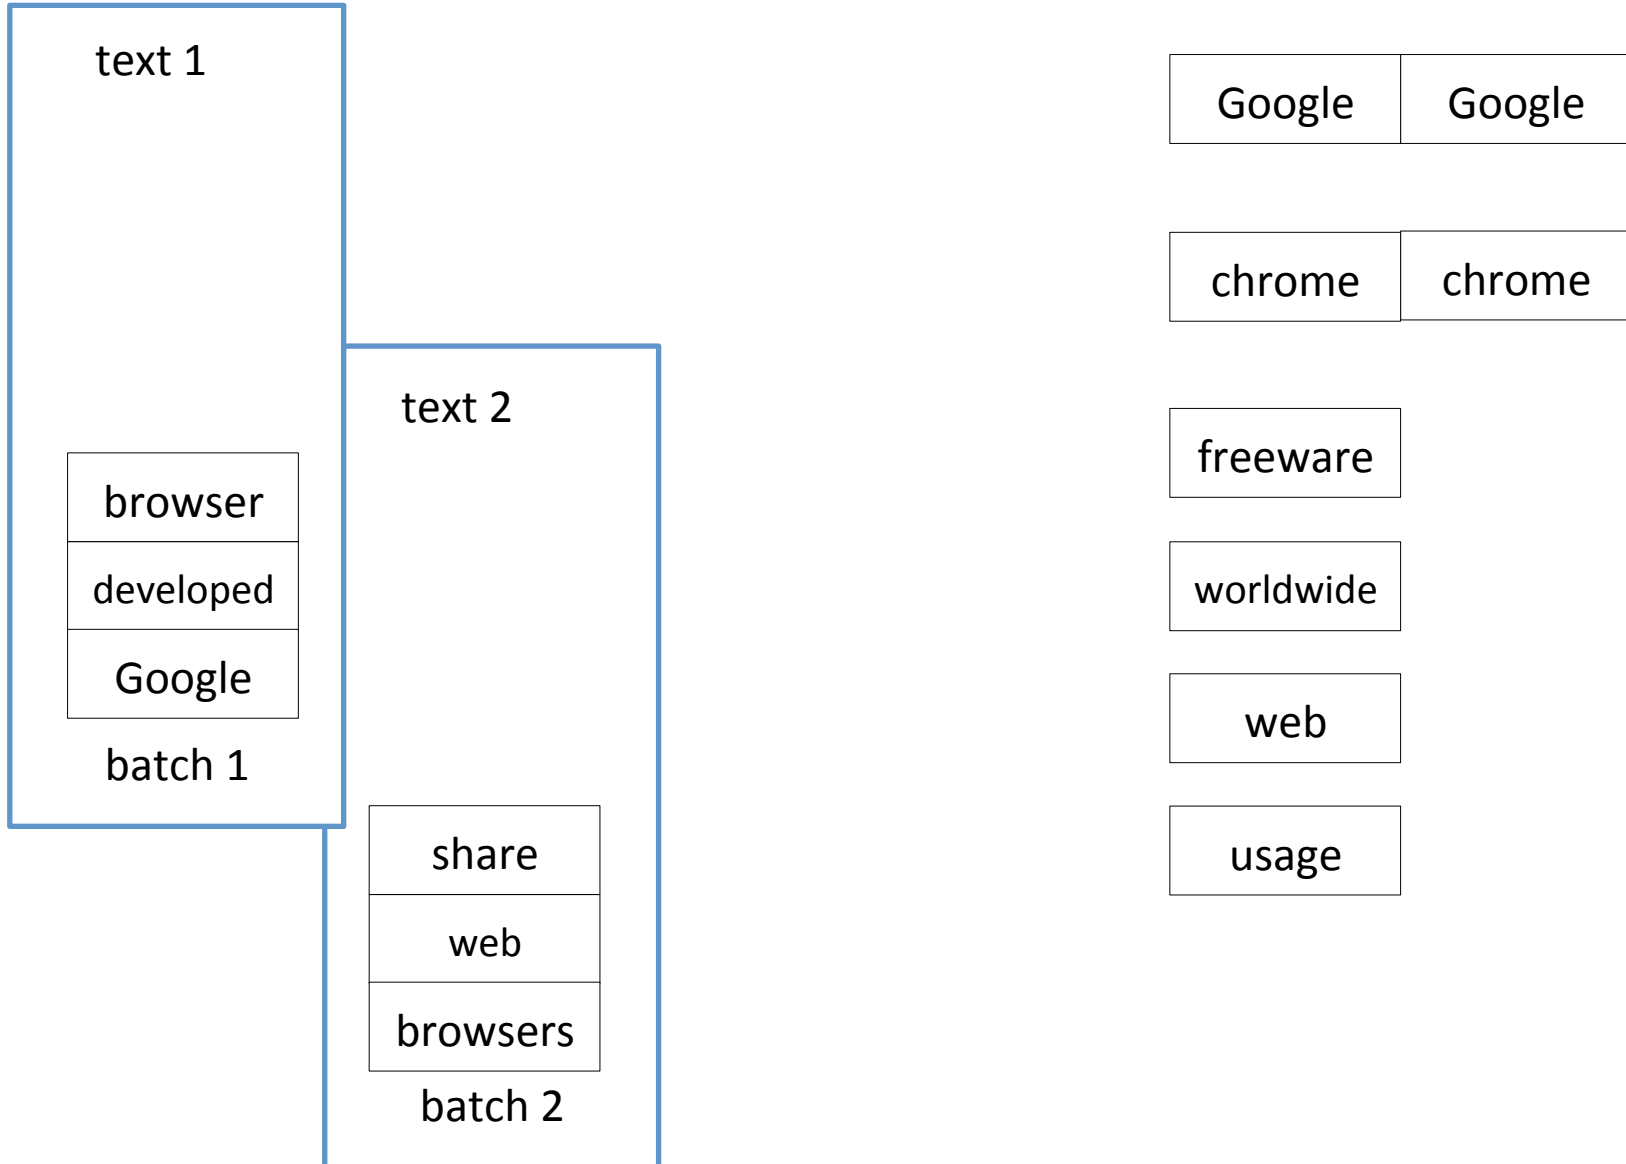




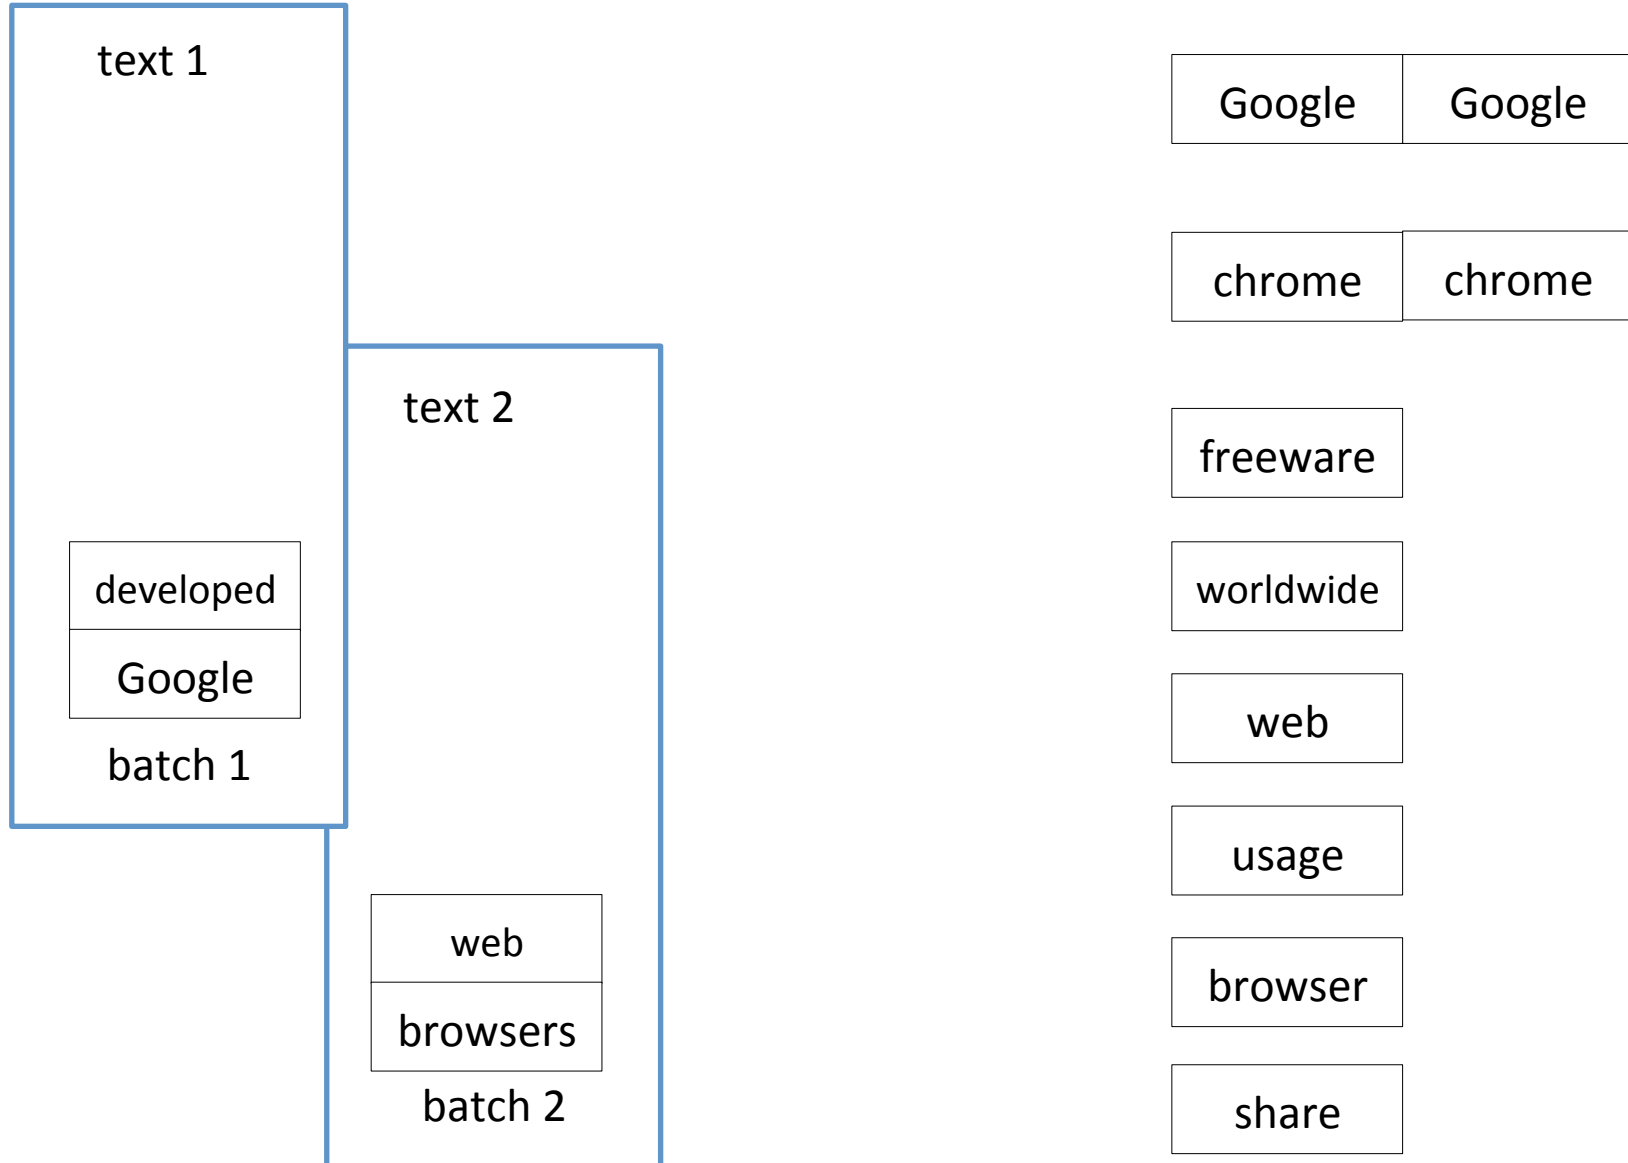
# M/R Computing Model



# M/R Computing Model




# M/R Computing Model



# M/R Computing Model



# M/R Computing Model



Google	Google	Google
--------	--------	--------

chrome	chrome
--------	--------

freeware
----------

worldwide
-----------

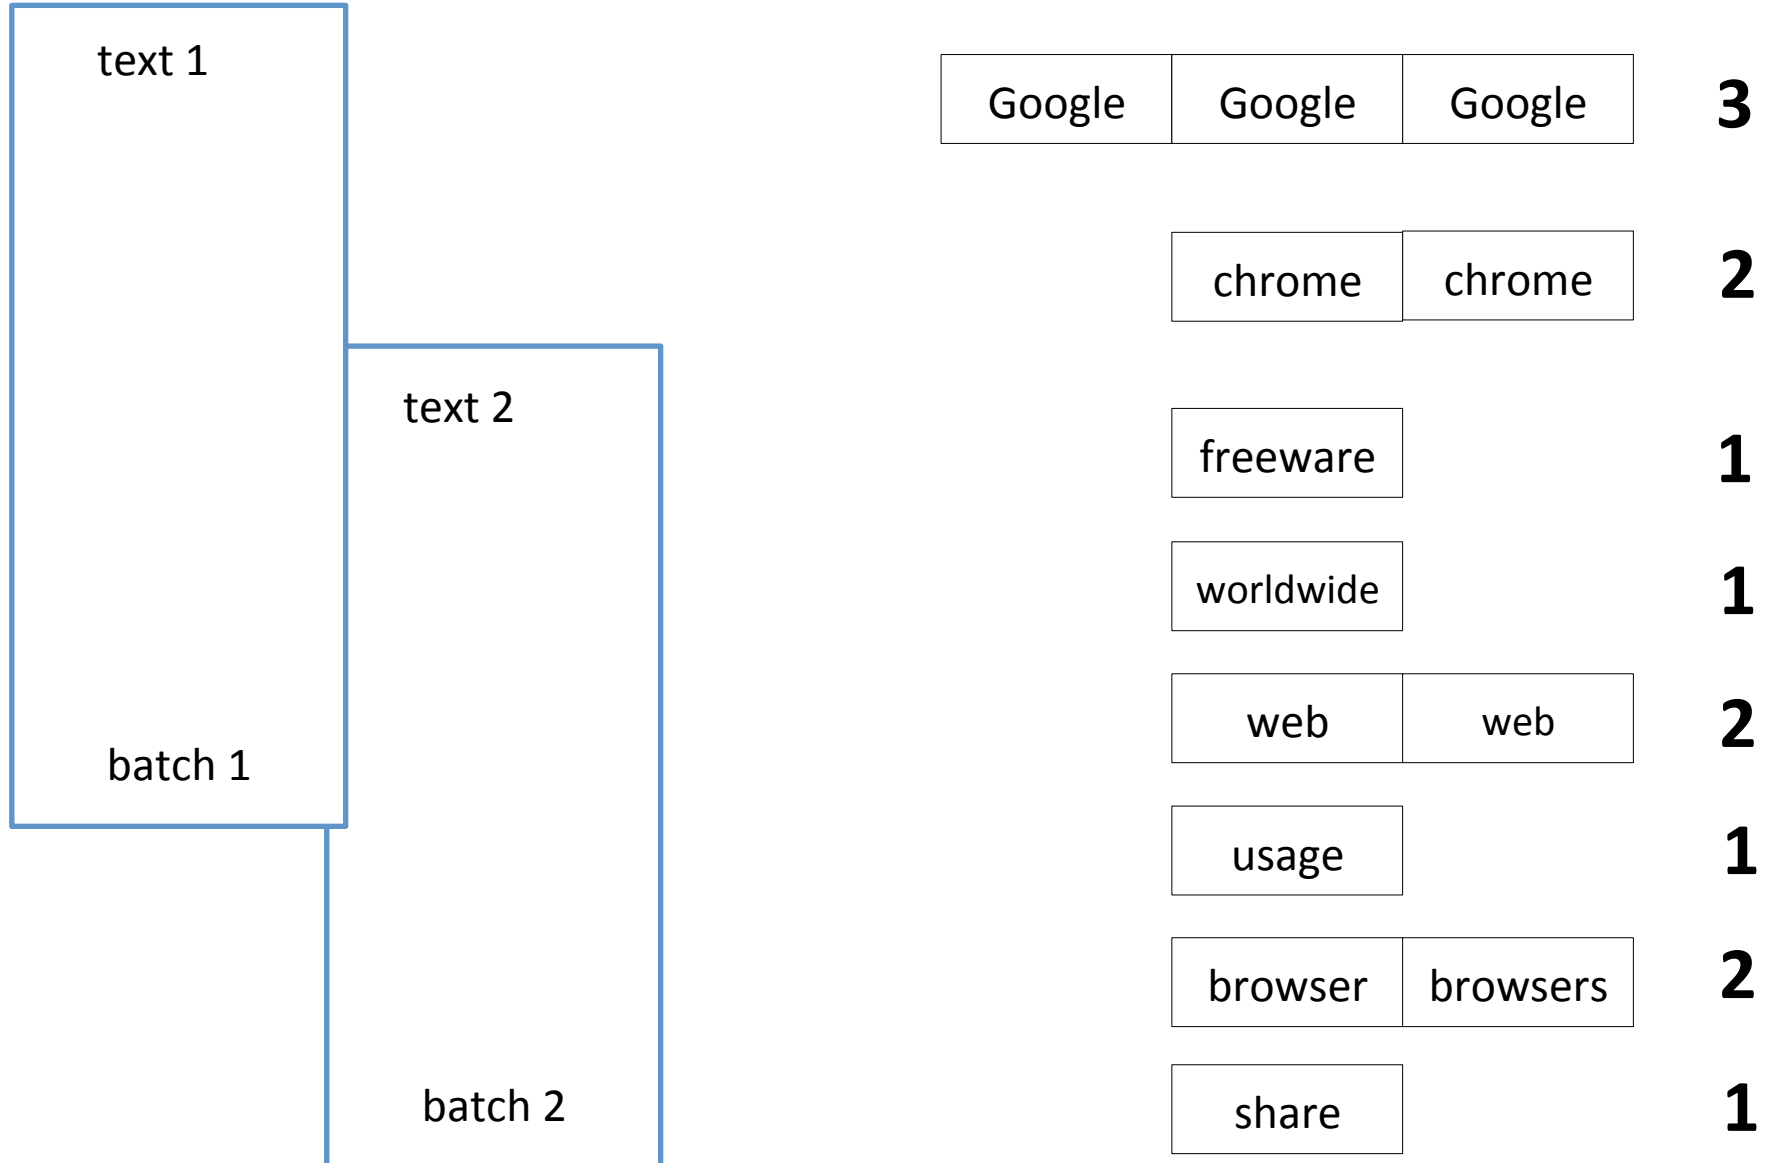
web	web
-----	-----

usage
-------

browser	browsers
---------	----------

share
-------

# M/R Computing Model




# Analyze search logs to find popular trends

## Google processes 3.5 billion searches per day

### Midterm Elections 2018

Hundreds of candidates vied for your vote across the US. See the top issues in search.




Search interest in voting, 10/30 to 11/06

[READ MORE](#) →

### Thanksgiving 2018

Thanksgiving falls on the 4th Thursday of November every year.



Most searched pies, past week US

[READ MORE](#) →

# Query Log Example

How many days until Thanksgiving?

What restaurants are open on Thanksgiving?

Is Trump party going to win the election ?

Where is Trump right now?

When is Thanksgiving?

Is Trump going to California ?

Can Trump win the next Presidency?

Why do we celebrate Thanksgiving?

When was the first Thanksgiving?

Why should I vote for Donald Trump ?

Why do people like Trump ?

What did Trump say ?

What tweeted Trump ?



# Split the log

How many days until Thanksgiving?

What restaurants are open on Thanksgiving?

Is Trump party going to win the election ?

Where is Trump right now?

When is Thanksgiving?

Is Trump going to California ?

Can Trump win the next Presidency?

-----

Why do we celebrate Thanksgiving?

When was the first Thanksgiving?

Why should I vote for Donald Trump ?

Why do people like Trump ?

What did Trump say ?

What tweeted Trump ?

# Then analyze

How many days until Thanksgiving?

What restaurants are open on Thanksgiving?

Is Trump party going to win the election ?

Where is Trump right now?

When is Thanksgiving?

Thanksgiving	5
--------------	---

Is Trump going to California ?

Can Trump win the next Presidency?

Trump	8
-------	---

-----

Why do we celebrate Thanksgiving?

When was the first Thanksgiving?

Why should I vote for Donald Trump ?

Why do people like Trump ?

What did Trump say ?

What tweeted Trump ?

# Query Processing

- Hadoop-M/R is **not** a data management system, it is a general framework.
- It can therefore implement queries :
  - *does not have better performances than a DW*
  - *but easier to setup & run, and more flexible*




# Group By




```
SELECT store_id, sum(sale_amount)
FROM sales
GROUP BY store_id
```



# Group-by in M/R

store\_id sale\_amount

store\_id sale\_amount



	10
	30
	20



	40
	50
	10

	80
	60

# Group-by in M/R



store\_id sale\_amount


	30
	20

	50
	10

	60
---	----

store\_id sale\_amount


	10
	40

	80
---	----




# Group-by in M/R


store\_id sale\_amount



	20
---	----

	10
--	----

store\_id sale\_amount

	10
	40
	60




	30
---	----




	80
	50



# Group-by in M/R

store\_id sale\_amount

store\_id sale\_amount

	10
	40
	60

	30
	20
	10

	80
	50






# Group-by in M/R



**store\_id sale\_amount**

**store\_id sale\_amount**

SUM(sale\_amount)  
= 110



	30
	20
	10

	80
	50

# Group-by in M/R

**store\_id sale\_amount**

**store\_id sale\_amount**

SUM(sale\_amount)  
= 110



SUM(sale\_amount)  
= 60



SUM(sale\_amount)  
= 130









# Join

```
SELECT  store_name, sale_amount
FROM    sales, store
WHERE
sales.store_id = store.store_id
```

# Join in M/R

**store\_id**  
**sale\_amount**



	10
	30
	20
	40
	50
	10



**store\_id**   **name**

	Green
	Blue
	Red



# Join in M/R



store\_id  
sale\_amount


	30
	20

	50
	10



store\_id   name


	Blue
	Red



	10
	40


	Green
---	-------


# Join in M/R


	10
	40

	Green
---	-------



	30
	20


	Blue
---	------

	50
---	----

	Red
---	-----



# Join in M/R


	10
	40

	Green
---	-------

(10, Green)


(40, Green)


	30
	20

	Blue
---	------

(30, Blue)

(20, Blue)


	50
---	----

	Red
---	-----

(50, Red)

# Social Network Analysis

- Find all pairs of "similar" users
  - in terms of interests, age, country, behavior ...




- Worst-case  $n(n-1)/2$  comparisons ( $n=\text{\#users}$ )




# User-Similarity in M/R

**user**





# User-Similarity in M/R

**user**




# User-Similarity in M/R

user



$$n(n-1)/2$$

# User-Similarity in M/R



# It is called M/R but...

- as we have seen it is rather :
  - Map
  - Shuffle (Regroup)
  - Reduce
- why is that ?
  - because the programmer just writes the Map and Reduce functions !

# **MAP-REDUCE PROGRAMMING**

# Central Notion

In MR very operation is expressed by using pairs

`<key, value>`

Keys are not only numbers !

Values are not only text !

# Map(key $k_{\text{input}}$ , value $v$ ) $\rightarrow$ Set<key,value>

Read : map function takes in input a key-value pair and outputs a set of (key,value) pairs

A Map-call is executed for every ( $k_{\text{input}}$  ,  $v$  ) pair

- Word count :  $k_{\text{input}}$  is a line-id  $v$  is a line of text
- Trends :  $k_{\text{input}}$  is a line-id  $v$  is a log line
- Group by :  $k_{\text{input}}$  is a tuple id  $v$  is a tuple
- Join :  $k_{\text{input}}$  is a tuple id  $v$  is a tuple
- Similarity :  $k_{\text{input}}$  is a user id  $v$  is a user



**Map(key  $k_{\text{input}}$ , value  $v$ )  $\rightarrow$  Set<key,value>**

Read : map function takes in input a key-value pair  
and outputs a set of (  $k_{\text{group}}$  ,  $v$  ) pairs

A Map-call is executed for every (  $k_{\text{input}}$  ,  $v$  ) pair

- Word count :  $k_{\text{input}}$  is a line-id  $v$  is a line of text
- Trends :  $k_{\text{input}}$  is a line-id  $v$  is a log line
- Group by :  $k_{\text{input}}$  is a tuple id  $v$  is a tuple
- Join :  $k_{\text{input}}$  is a tuple id  $v$  is a tuple
- Similarity :  $k_{\text{input}}$  is a user id  $v$  is a user

# Reduce( key $k_{\text{group}}$ , Set<value> $V$ ) → Set<key,value>

Read : reduce function takes in input a key and a set of values and outputs a set of key-value pairs

All ( $k_{\text{group}}$ ,  $V$ )-pairs for a given  $k_{\text{group}}$  are evaluated by the same reducer !

Wordcount :             $k_{\text{group}}$  is a word             $V$  number of occurrences

Trends :             $k_{\text{group}}$  is a word             $V$  number of occurrences

Group by :             $k_{\text{group}}$  is a group-by attribute-value     $V$  a set of tuples

Join :             $k_{\text{group}}$  is a join attribute-value             $V$  a set of tuples

Similarity :             $k_{\text{group}}$  is a user-user pair id             $V$  two users' profiles

# Keyword Count Using MapReduce

```
map(key k, value phrase):
```

```
    for each word in phrase :
```

```
        emit( word , 1 )    //generates a <key,value> pair
```

```
reduce(key word, values occurrences):
```

```
    emit( key , occurrences.size() )
```

# Word Count

phrase 1

Google
chrome
freeware
web
browser
developed
Google

phrase 2

Google
chrome
worldwide
usage
share
web
browser

Google,1	Google,1	Google,1
----------	----------	----------

**3**

chrome,1	chrome,1
----------	----------

**2**

freeware,1
------------

**1**

worldwide,1
-------------

**1**

web,1	web,1
-------	-------

**2**

usage,1
---------

**1**

browser,1	browser,1
-----------	-----------

**2**

share, 1
----------

**1**

# Query Trends Using MapReduce

```
map(key k, value batch_of_lines):
```

```
    for each word in batch_of_lines:  
        emit( word , 1 )
```

```
reduce(key word , values occurrences):
```

```
    emit( key, occurrences.size() )
```

# Group-by Using MapReduce

```
map(key k, value tuples):
```

```
    for each tuple in tuples:
```

```
        emit( tuple.store_id , tuple.sale_amount )
```

```
reduce(key store_id, values sales):
```

```
    total_sales=0;
```




```
    for each s in sales
```




```
        total_sales+=s;
```



```
    emit( store_id , total_sales )
```

# Group By in M/R

store\_id sale\_amount

	10
	30
	20

	40
	50
	10

	80
	60

store\_id sale\_amount

SUM(sale\_amount)  
= 100



30
20
10



80
50



# Join Using MapReduce

**map(key k, value tuples):**

```
for each tuple t in tuples:
    if t belongs to SALES
        emit( t.store_id , <"profit", t.sale_amount> )
    if t belongs to STORES
        emit( t.store_id , <"store", t.store_name> )
```







**reduce(key store\_id, values mixed-attributes):**



```
for each a in mixed-attributes
    for each b in mixed-attributes
        if a[1]=="sale" and b[1]=="store"
            emit( store_id , <a[2],b[2]> )
```




# Join in M/R

**store\_id**  
**sale\_amount**


	10
	30
	20
	40
	50
	10

<b>store_id</b>	<b>name</b>
	Green
	Blue
	Red

	sale_amount	10	store	Green
	sale_amount	40		

(10, Green)

(40, Green)

	sale_amount	30	store	Blue
	sale_amount	20		

(30, Blue)

(20, Blue)

	sale_amount	50	store	Red

(50, Red)

# User Similarity Using MapReduce

```
map(key user_i_id, value user_i_profile):
```

```
// send the profile of user i to all other users
```

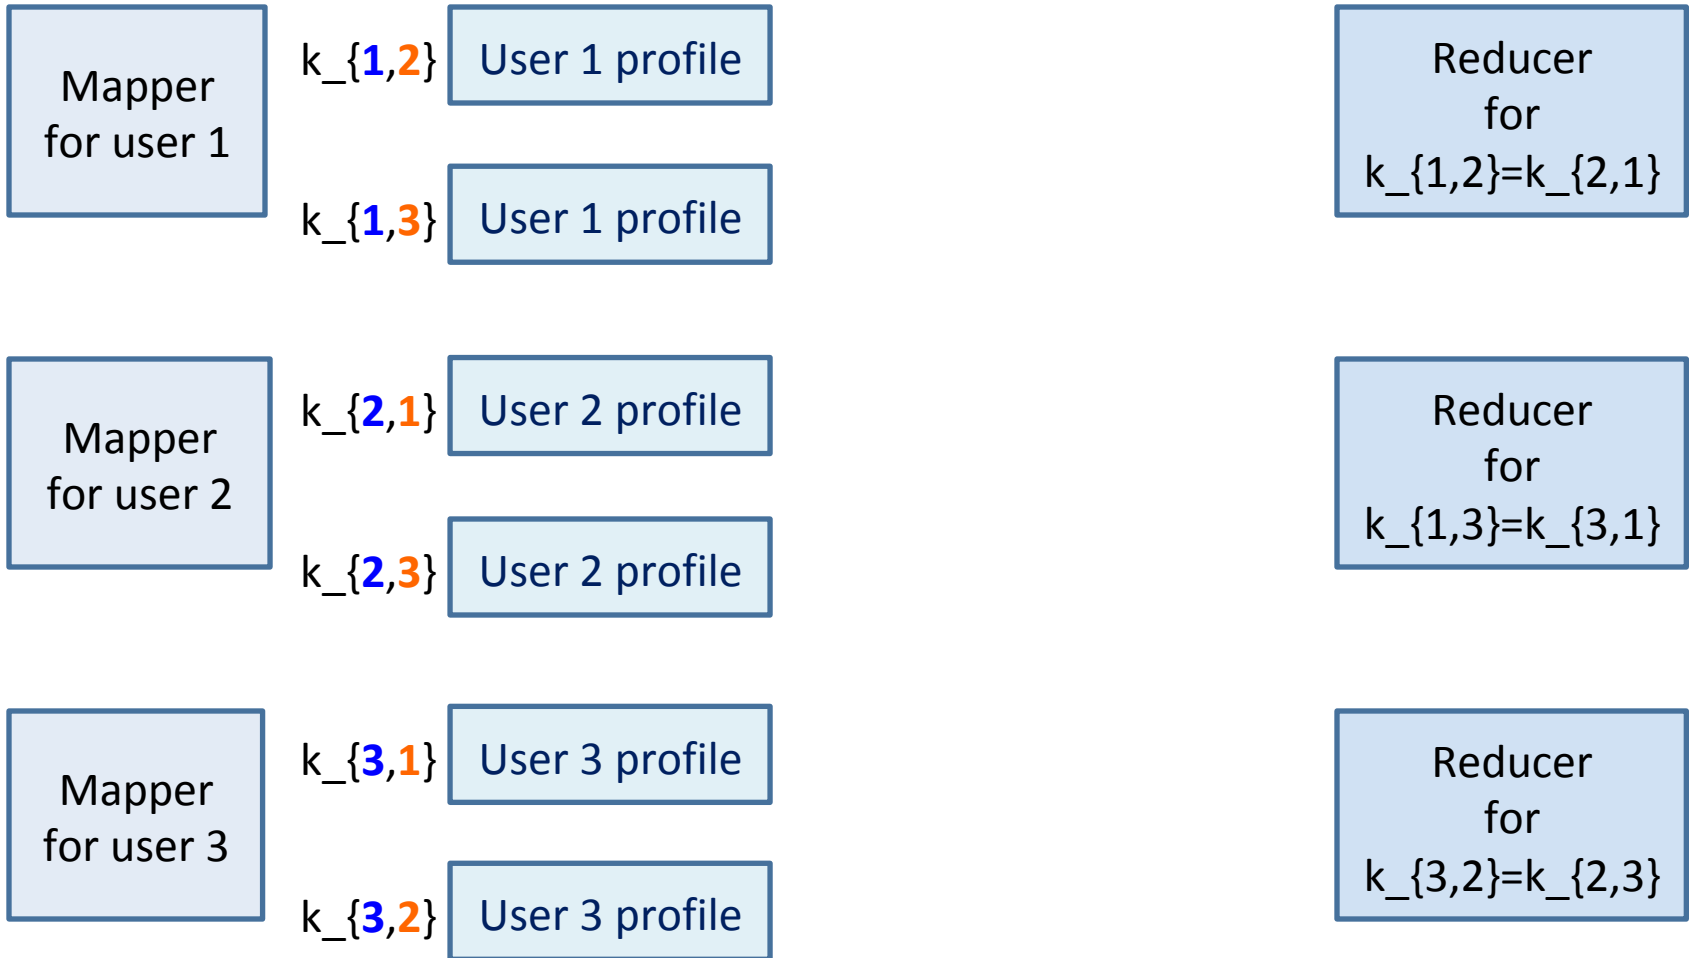
```
reduce(key k_ij , values two_user_records):
```

```
// compare two users
```


# User Similarity Using MapReduce

```
map(key user_i, value user_i_profile):  
    for each user_j  
        create a "new" destination-key k_{i,j}  
        emit( k_{i,j} , user_i_profile )
```

# Example: Three Users



# Example: Three Users



# User Similarity Using MapReduce

```
map(key user_i, value user_i_profile):  
    for each user_j                                with i != j  
        create a key k_{i,j}                      with k_{i,j}=k_{j,i}  
        emit( k_{i,j} , user_i_profile)
```

# User Similarity Using MapReduce

```
map(key user_i, value user_i_profile):  
    for each user_j                                with i != j  
        create a key  $k_{\{i,j\}}$                     with  $k_{\{i,j\}}=k_{\{j,i\}}$   
        emit(  $k_{\{i,j\}}$  , user_i_profile)
```

```
reduce(key  $k_{ij}$  , values two_user_records):
```

```
    u1 = two_user_records[1]
```

```
    u2 = two_user_records[2]
```

```
    if similarity( u1 , u2 ) >= 0.9
```

```
        emit( "similar" , < u1.id , u2.id > )
```

# Map-Reduce Programming

**There is not a single line of code dedicated to parallelization !!**

## **Map-Reduce environment takes care of:**

- Partitioning the input data
- Scheduling the program's execution across a set of machines
- Performing the **group by key/shuffle** step
- Handling machine failures
- Managing required inter-machine communication



Jeffrey Ullman

# **WHAT CAN GO WRONG**

# The “Drug Interaction” Problem

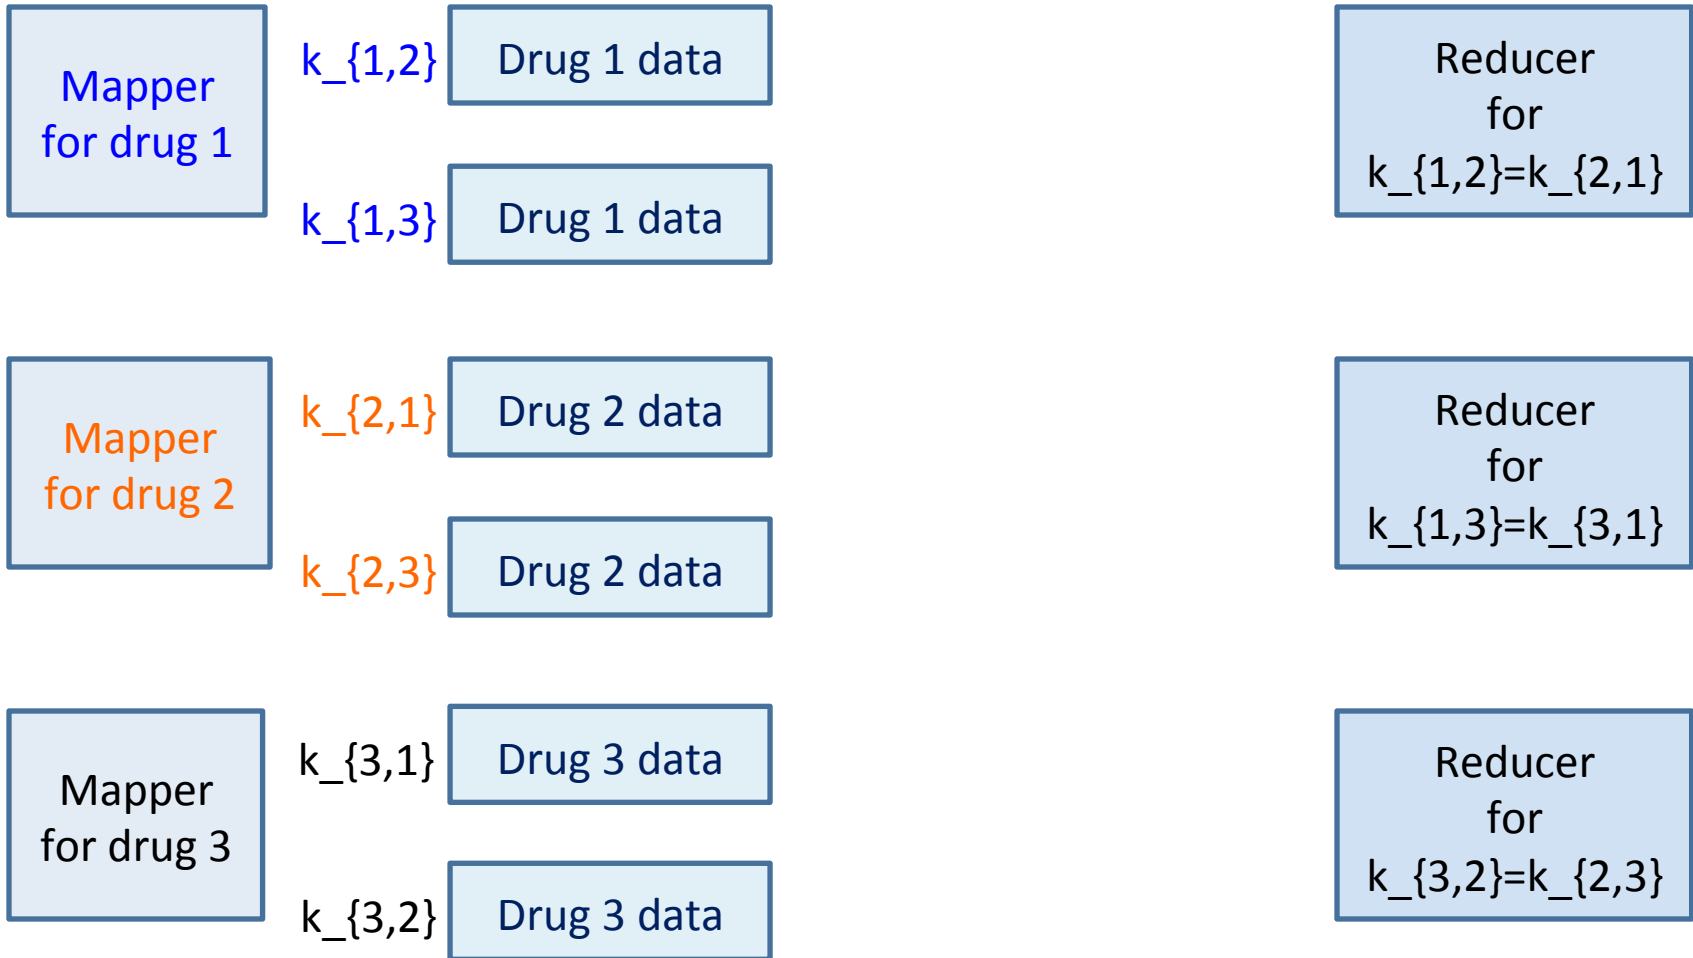
- Data consists of records for 3000 drugs.
  - List of patients taking them, dates, diagnoses.
  - About 1M of data per drug.
- Problem is to find drug interactions.
  - **Example**: two drugs that when taken together increase the risk of heart attack.
- Must examine each pair of drugs and compare their data using statistical tests.

# “Drug Interaction” Using MapReduce

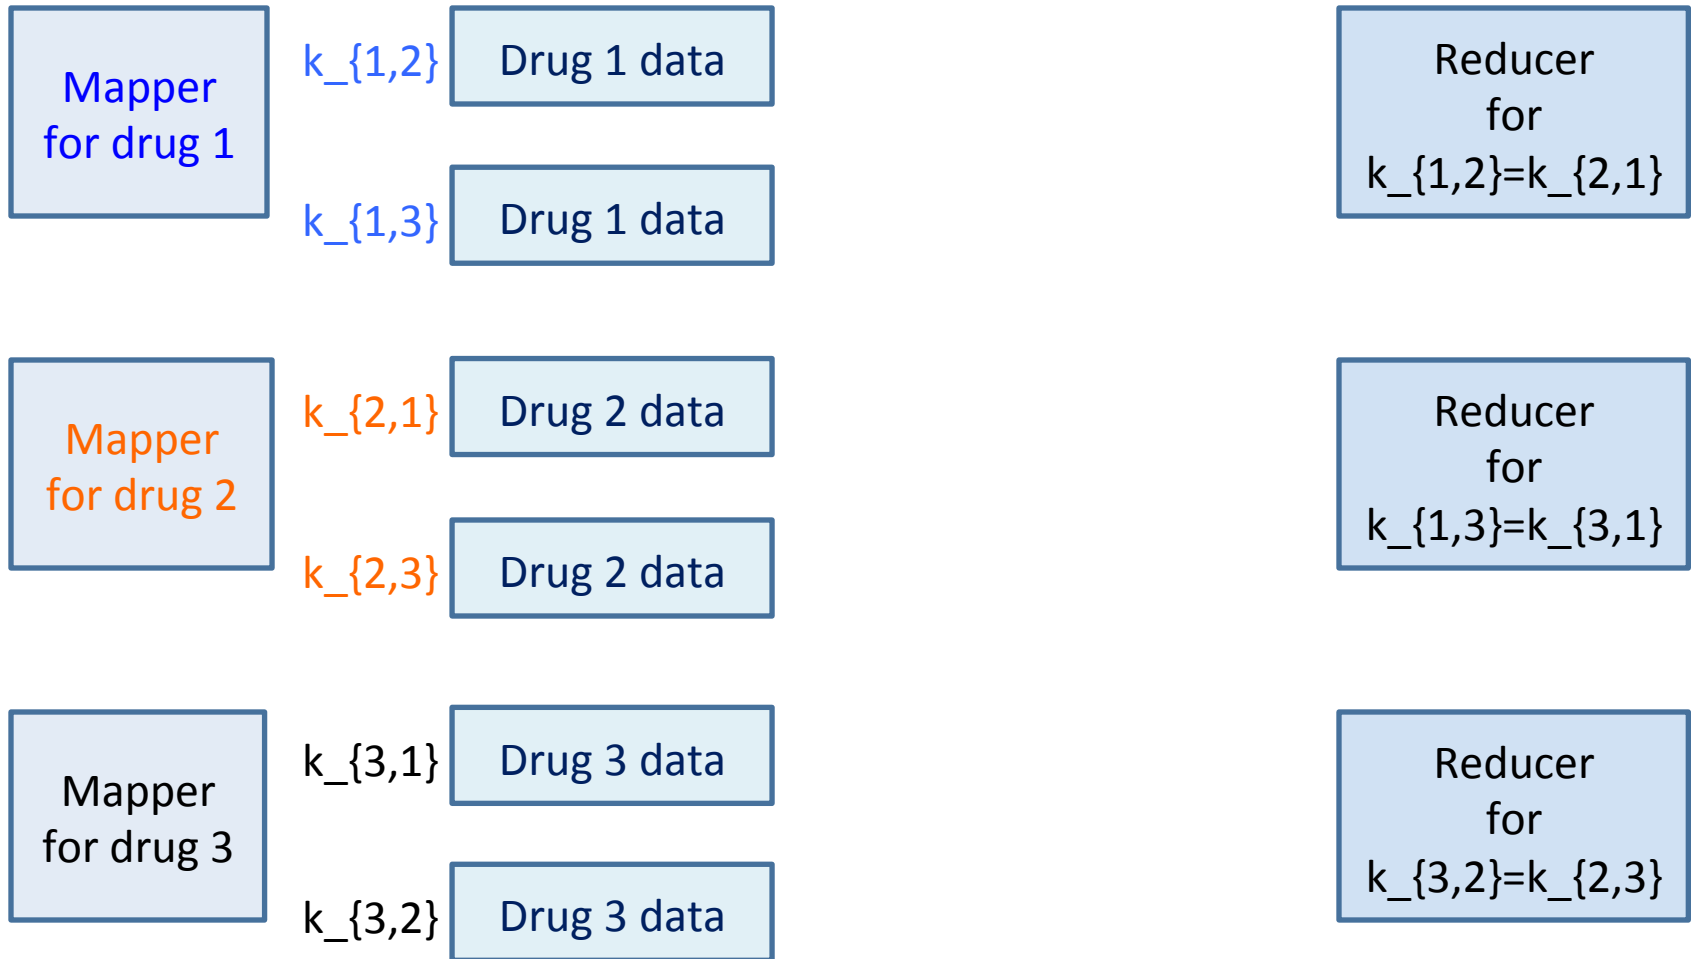
```
map(key drug_i, value drug_i_record):  
    for each j in 1..3000 with i != j  
        create a key  $k_{\{i,j\}}$  such that  $k_{\{i,j\}}=k_{\{j,i\}}$   
        emit(  $k_{\{i,j\}}$  , drug_i_record )
```

```
reduce(key k_drug_pair , values two_drug_records):  
  
    d1=two_drug_records [1]  
    d2=two_drug_records [2]  
  
    if statistical-test-significative( d1 , d2 )  
        emit( "interacting" , < d1.id , d2.id > )  
    else  
        emit( "non-interacting" , < d1.id , d2.id > )
```

# Example: Three Drugs



# Example: Three Drugs



# What Went Wrong?

- 3000 drugs → 3000 map tasks
- each sends 2999 copies of a single drug record
- which amounts to 1MB
- = 9TB communicated over a 1Gb Ethernet
- ~ 90,000 seconds (25h) of network use.
  - assuming no other job is using the network

# A Better Approach

- The way to handle this problem is to group the drugs
- For example : 30 groups of 100 drugs each
- This way, a single drug record is replicated 29 times instead of 2999

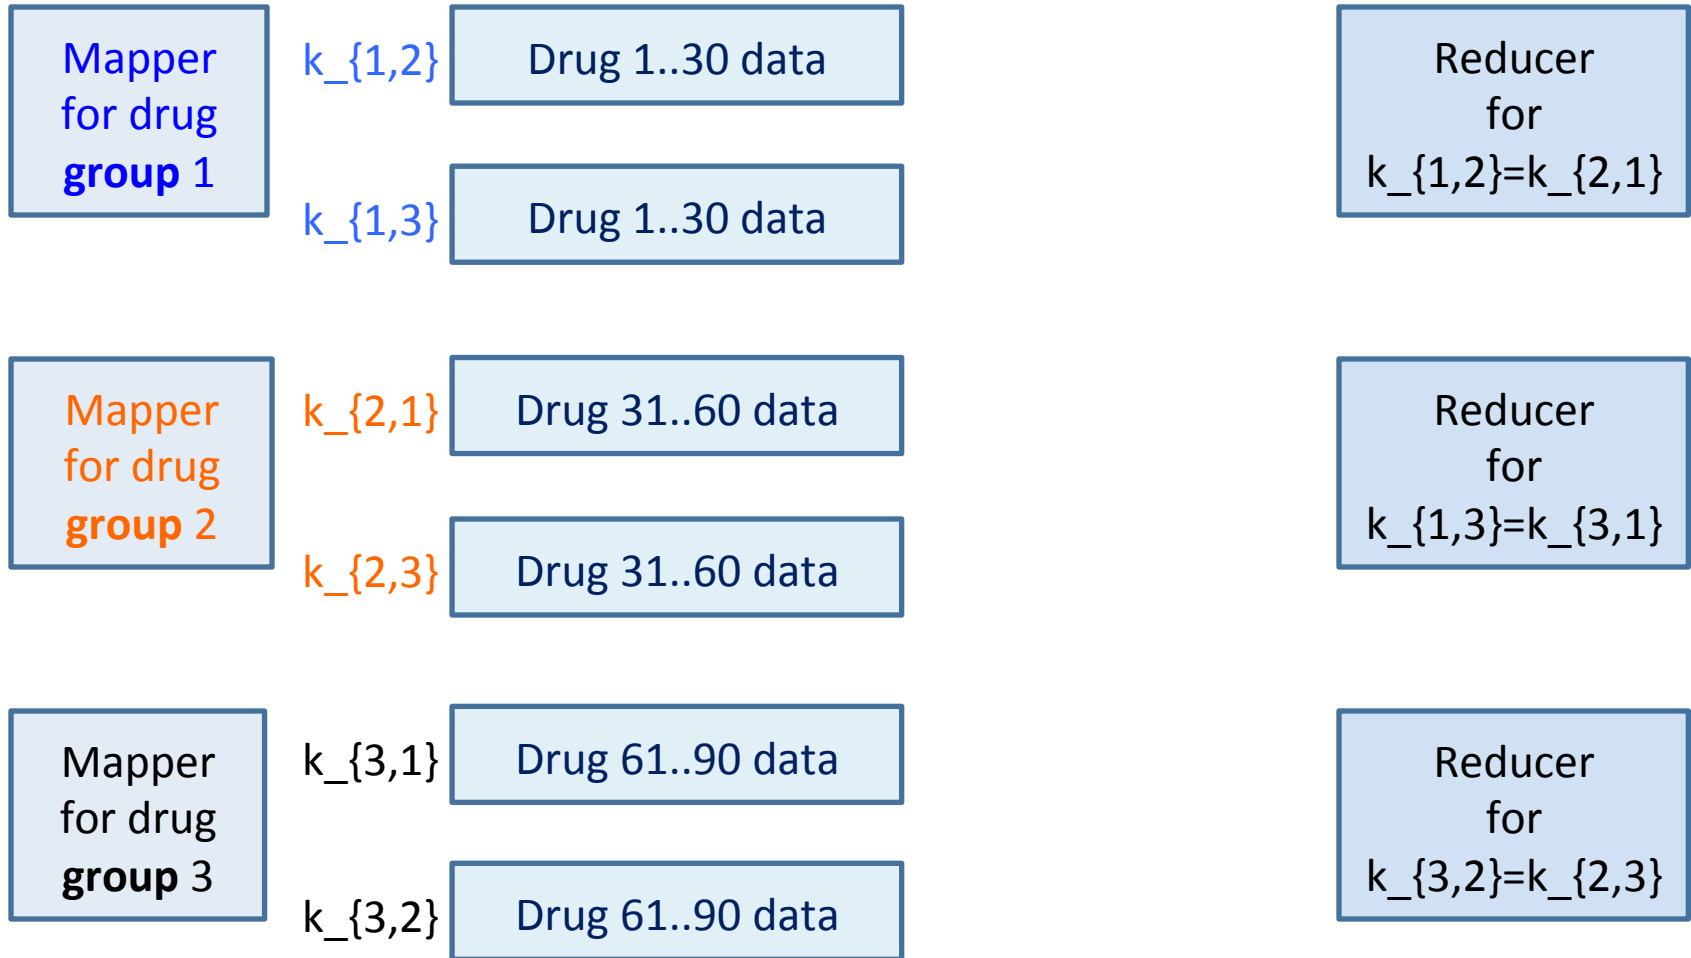
# Drug Interaction Using MapReduce (2)

```
map(key drug_group_i_id, value drug_group_i_record) :  
    for each j in 1..30 with i != j  
        create a key  $k_{\{i,j\}}$  such that  $k_{\{i,j\}}=k_{\{j,i\}}$   
        emit(  $k_{\{i,j\}}$  , drug_group_i_record )
```

```
reduce(key k_ij , values two_groups_records) :  
  
    g1=two_groups_records [1]  
    g2=two_groups_records [2]  
    for each d1 in g1  
        for each d2 in g2  
            if statistical-test-significative( d1 , d2 )  
                emit( "interacting" , < d1.id , d2.id > )  
            else  
                emit( "non-interacting" , < d1.id , d2.id > )88
```




# Example: Three Drugs



# Why It Works

- The big difference is in the communication requirement.
- Now, each of 3000 drugs' 1MB records is replicated 29 times.
  - Communication cost = 87GB, vs. 9TB.



# Cost Measures for Algorithms

- In MapReduce we quantify the cost of an algorithm using
  1. *Communication cost* = total I/O of all processes
  2. *Computation cost* = total CPU time of all processes

# Why is this important ?

- On a **public cloud**, you **pay for computation** and you also **pay for communication**.
  - Balancing the two is an important part of algorithm design.
- **If communication cost dominates total cost** it influences how much parallelism you can extract from an algorithm.
  - time reductions are not as good as expected

# Reducer size is a key point

- In many cases, the big issue is whether a reducer has too much input to operate in main memory.
  - To get reducers with small input size, you may need a lot of communication.
- The “Drug-Interaction Problem” is a good model for how one can trade off communication against parallelism.

# And Why User Similarity Works ?

- 3000 users → 3000 map tasks
- each sends 2999 copies of a each user record
- which amounts to ~1KB
- = 9 GB communicated over a 1Gb Ethernet
- ~ 90 seconds of network use.

**ARCHITECTURE**




# Big-Data Cluster

Today, a standard architecture for Big-Data is emerging:

- Cluster of commodity Linux nodes
- Commodity network (ethernet) to connect them
- Nodes Organized into racks
  - Intra-rack connection typically gigabit speed.
  - Inter-rack connection faster by a small factor.
- Shared-nothing : no shared memory
  - this is not High Performance Computing (HPC)

# Cluster Architecture



Each rack contains 16-64 nodes



In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>



# Challenges With Implementing M/R

- Distributing computation over a network can be non-trivial
- Data is huge and copying data over a network takes time
  - we have just seen an example
- Machines fail:
  - server may stay up 3 years (1,000 days) ;  
with 1000 servers expect to loose 1/day

# Data Replication



Each rack contains 16-64 nodes

# Distributed File System

- Chunk Servers.
  - File is split into contiguous chunks, typically 64MB.
  - Each chunk replicated (usually 2x or 3x).
  - Try to keep replicas in different racks.
- Master Node for a file.
  - Stores metadata, location of all chunks.
  - Possibly replicated.

# Distributed File System

## Client library for file access

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data
- **Try to send map computation where the data is**
  - **but cannot send too many jobs to the same machines, data could be moved before map is executed**
- During shuffle send all key-value pairs to the same reduce machine
  - better if closed to where the map has been done
    - but this is not always possible